Training Normalizing Flows

Maximum likelihood training to find parameters θ that make our training samples likely under our parameterized model.

$$\max_{\theta} \mathbb{E}_{\mathbf{x} \sim p_x} p_{\theta}(\mathbf{x}) = \min_{\theta} \mathbb{E}_{\mathbf{x} \sim p_x} - \log p_{\theta}(\mathbf{x})$$

Training Normalizing Flows

Maximum likelihood training to find parameters θ that make our training samples likely under our parameterized model.

$$\max_{\theta} \mathbb{E}_{\mathbf{x} \sim p_x} p_\theta(\mathbf{x}) = \min_{\theta} \ \mathbb{E}_{\mathbf{x} \sim p_x} - \log p_\theta(\mathbf{x})$$

$$= \min_{\theta} \ \mathbb{E}_{\mathbf{x} \sim p_x} - \log \left[p_Z(T_\theta(\mathbf{x})) \left| \det \frac{\partial T_\theta(\mathbf{x})}{\partial \mathbf{x}} \right| \right]$$