

Training Normalizing Flows

Maximum likelihood training to find parameters θ that make our training samples likely under our parameterized model.

$$\max_{\theta} \mathbb{E}_{\mathbf{x} \sim p_x} p_{\theta}(\mathbf{x})$$

Training Normalizing Flows

Maximum likelihood training to find parameters θ that make our training samples likely under our parameterized model.

$$\max_{\theta} \mathbb{E}_{\mathbf{x} \sim p_x} p_{\theta}(\mathbf{x}) = \min_{\theta} \mathbb{E}_{\mathbf{x} \sim p_x} -\log p_{\theta}(\mathbf{x})$$